

PROBLEM STATEMENT:

IBM faces a critical challenge with employee attrition. Understanding its drivers—job satisfaction, work-life balance, compensation, career growth, and organizational culture—is paramount. Through data analytics and predictive modelling, IBM aims to forecast attrition risks and implement targeted retention strategies. These include enhancing engagement initiatives, refining performance management, revising compensation, fostering a positive work environment, and improving leadership. Continuous evaluation ensures effectiveness, enabling iterative improvements. By addressing these factors, IBM endeavors to reduce attrition, nurture a stable, motivated workforce, and sustain its competitive edge.

INTRODUCTION:

- Using machine algorithms to predict employee attrition offers a comprehensive scope, involving data analysis, prediction modelling, and risk identification. By analysing historical data, these algorithms identify patterns and influential factors contributing to attrition, enabling organizations to predict future turnover rates and identify high-risk employees. This insight allows for proactive intervention and targeted retention strategies to retain valuable talent. Continuous refinement of predictive models ensures ongoing effectiveness in mitigating attrition and fostering a stable workforce.

OBJECTIVES:

Data Acquisition:

- We , likely won't have access to real IBM employee data due to privacy concerns. However, we can leverage publicly available datasets like the IBM HR Attrition dataset on Kaggle <https://www.kaggle.com/datasets/pavansubhasht/ibm-hr-analytics-attrition-dataset>.

Data Splitting and Preparation:

- The training set will be used to build the model, and the testing set will be used to evaluate its accuracy in predicting employee attrition.
- Explore and understand the features in the dataset. These might include factors like:
- Employee demographics (age, gender, department)
- Job details (job satisfaction, work-life balance, years with company)
- Performance metrics (salary, promotions, trainings received)

Data Preprocessing:

- Handle missing values, outliers, and any data inconsistencies through appropriate preprocessing techniques to ensure the dataset's quality.

Machine Learning Model Building:

- Apply various machine learning algorithms typically used for classification tasks, such as:
- Logistic Regression: A classic algorithm for predicting binary outcomes (employee leaving or staying)
- Decision Trees: Easy to interpret and understand, providing insights into key factors influencing attrition.

Evaluation and Feature Engineering:

- Train each model on the training set and evaluate its performance on the testing set using metrics like accuracy, precision, and recall.

Learning Outcomes:

- Gain experience with Python libraries like NumPy and scikit-learn for data manipulation and machine learning model development.
- Understand the process of building and evaluating machine learning models for real-world business problems like employee retention.

Solution Approach:

- **Logistic Regression:**

Predicts the probability of an employee leaving based on their features, allowing for classification into high or low-risk categories. Analysing the model's coefficients will reveal which factors have the strongest positive or negative influence on employee retention.

- **K-Nearest Neighbors (KNN):**

Classifies employees based on the similarity of their features to those who previously left IBM. Here, choosing the optimal K value and potentially scaling features are crucial.

- **Decision Tree:**

Creates a tree-like structure that recursively splits the data based on the most important factors influencing employee departure. This method offers valuable insights into the key drivers of attrition.

- **Gradient Boosting Classifier:**

Builds an ensemble of decision trees, where each tree focuses on correcting the errors of the previous one. This ensemble approach can lead to a robust and accurate model.

- **Random Forest:**

Similar to gradient boosting, it builds an ensemble of decision trees but introduces randomness in feature selection at each split. This helps prevent overfitting and improve generalizability across different departments or demographics

- **Support Vector Machine (SVM):**

Working: SVMs aim to find a hyperplane in the feature space that best separates the data points representing employees who left

(positive class) from those who stayed (negative class). This hyperplane maximizes the margin between the classes, leading to a robust decision boundary.

Considerations: Choosing the right kernel function (e.g., linear, radial basis) is crucial for SVM performance in employee attrition prediction. Feature scaling might also be necessary.

- **Neural Network:**

Working: Neural networks are inspired by the human brain and consist of interconnected layers of artificial neurons. These neurons learn complex patterns from the data to predict employee attrition.

Considerations: Neural networks can be powerful but require careful tuning of hyperparameters (e.g., number of layers, neurons per layer) to avoid overfitting and achieve optimal performance.

- **XGBoost (Extreme Gradient Boosting):**

Working: XGBoost is an ensemble learning technique that builds a series of decision trees sequentially. Each tree focuses on correcting the errors of the previous one, resulting in a highly accurate model for predicting employee attrition. Benefits: XGBoost offers built-in regularization to prevent overfitting and handles missing values effectively. It can also be interpretable to some extent, providing insights into the factors influencing employee departure.

Timeline:

- **Phase 1: Planning (Days 1-3)**

- Day 1: Define project goals and scope (focus on predicting employee attrition).
- Day 2: Identify relevant employee data (features) and target variable (employee leaving or staying).
- Day 3: Develop a basic workflow for data processing, model building, and evaluation.

Phase 2: Design (Days 4-7)

- Day 4: Download and explore the chosen employee dataset.
- Day 5: Analyze data distribution (e.g., histograms, boxplots) to identify potential issues like missing values or outliers.
- Day 6: Plan for data cleaning and pre-processing steps.
- Day 7: Choose machine learning algorithms to evaluate (e.g., Logistic Regression, Random Forest, XGBoost).

Phase 3: Develop (Days 8-10)

- Model Training (Days 8-9):
- Day 8: Pre-process data (handle missing values, categorical encoding, feature scaling if necessary).
- Day 9: Split data into training and testing sets. Train various machine learning models on the training set.
- Model Testing (Day 10):
- Evaluate model performance using metrics like accuracy, precision, recall, and F1-score on the testing set.
- Compare model performance and choose the best performing model for attrition prediction.

Phase 4: System Enhancement, Deployment, Release (Days 11-15)

- Enhancement (Days 11-12):
- Day 11: Refine the chosen model based on evaluation results (e.g., hyperparameter tuning).
- Day 12: Implement feature engineering techniques (create new features) to potentially improve model performance.
- Deployment (Days 13-14):
- Day 13: Develop a basic Flask application to deploy the model in a controlled environment.
- Day 14: Perform end-to-end testing of the deployed application.

- Release (Day 15):
- Document the project and prepare user instructions.
- Release the deployed application to a limited group of users for initial feedback

Dataset Overview:

1. Age
2. Attrition
3. Business Travel
4. Daily Rate
5. Department
6. Distance From Home
7. Education
8. Education Field
9. Employee Count
11. Gender
12. Hourly Rate
13. Employee Number
14. Hourly Rate
15. Job Involvement
16. Job Level
17. Job Role
18. Job Satisfaction
19. Marital Status
20. Monthly Income
21. Monthly Rate
22. Num Companies Worked
23. Over18
24. Over Time
25. Percent Salary Hike
26. Performance Rating
27. Relationship Satisfaction
28. Standard Hours
29. Stock Option Level
30. Total Working Years
31. Training Times Last Year

- 32. Work Life Balance
- 33. Years At Company
- 34. Years In Current Role
- 35. YearsSinceLastPromotion
- 36. Years With CurrManager

1) Importing necessary libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.metrics import accuracy_score, classification_report
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
```

```

from sklearn.neural_network import MLPClassifier
from xgboost import XGBClassifier
from sklearn.metrics import confusion_matrix, precision_score, recall_score, f1_score
# Update the path accordingly

```

2) Loading the dataset

```

from google.colab import files
uploaded = files.upload()
df = pd.read_csv('WA_Fn-UseC_-HR-Employee-Attrition.csv')

```

3) Data Exploration

```

# Display basic information and the first few rows of the dataset
print(df.info())
print(df.head())

```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1470 entries, 0 to 1469
```

```
Data columns (total 35 columns):
```

#	Column	Non-Null Count	Dtype
0	Age	1470 non-null	int64
1	Attrition	1470 non-null	object
2	BusinessTravel	1470 non-null	object
3	DailyRate	1470 non-null	int64
4	Department	1470 non-null	object
5	DistanceFromHome	1470 non-null	int64
6	Education	1470 non-null	int64
7	EducationField	1470 non-null	object
8	EmployeeCount	1470 non-null	int64
9	EmployeeNumber	1470 non-null	int64
10	EnvironmentSatisfaction	1470 non-null	int64
11	Gender	1470 non-null	object
12	HourlyRate	1470 non-null	int64
13	JobInvolvement	1470 non-null	int64
14	JobLevel	1470 non-null	int64
15	JobRole	1470 non-null	object
16	JobSatisfaction	1470 non-null	int64
17	MaritalStatus	1470 non-null	object
18	MonthlyIncome	1470 non-null	int64
19	MonthlyRate	1470 non-null	int64
20	NumCompaniesWorked	1470 non-null	int64
21	Over18	1470 non-null	object
22	OverTime	1470 non-null	object
23	PercentSalaryHike	1470 non-null	int64
24	PerformanceRating	1470 non-null	int64
25	RelationshipSatisfaction	1470 non-null	int64
26	StandardHours	1470 non-null	int64
27	StockOptionLevel	1470 non-null	int64
28	TotalWorkingYears	1470 non-null	int64

29	TrainingTimesLastYear	1470	non-null	int64
30	WorkLifeBalance	1470	non-null	int64
31	YearsAtCompany	1470	non-null	int64
32	YearsInCurrentRole	1470	non-null	int64
33	YearsSinceLastPromotion	1470	non-null	int64
34	YearsWithCurrManager	1470	non-null	int64

dtypes: int64(26), object(9)

memory usage: 402.1+ KB

None

	Age	Attrition	BusinessTravel	DailyRate		Department
\						
0	41	Yes	Travel_Rarely	1102		Sales
1	49	No	Travel_Frequently	279	Research &	Development
2	37	Yes	Travel_Rarely	1373	Research &	Development
3	33	No	Travel_Frequently	1392	Research &	Development
4	27	No	Travel_Rarely	591	Research &	Development

	DistanceFromHome	Education	EducationField	EmployeeCount
EmployeeNumber \				
0	1	2	Life Sciences	1
1				
1	8	1	Life Sciences	1
2				
2	2	2	Other	1
4				
3	3	4	Life Sciences	1
5				
4	2	1	Medical	1
7				

	RelationshipSatisfaction	StandardHours	StockOptionLevel	\
0	...	1	80	0
1	...	4	80	1
2	...	2	80	0
3	...	3	80	0
4	...	4	80	1

	TotalWorkingYears	TrainingTimesLastYear	WorkLifeBalance
YearsAtCompany \			
0	8	0	1
6			
1	10	3	3
10			
2	7	3	3
0			
3	8	3	3
8			
4	6	3	3
2			

YearsInCurrentRole	YearsSinceLastPromotion	YearsWithCurrManager
--------------------	-------------------------	----------------------

0	4	0	5
1	7	1	7
2	0	0	0
3	7	3	0
4	2	2	2

[5 rows x 35 columns]

S"mmaíQ statistics roí →"mcíical rcat"ícs:

```
print(df.describe())
```

Age	DailyRate	DistanceFromHome	Education	EmployeeCount \
count	1470.000000	1470.000000	1470.000000	1470.000000
1470.0				
mean	36.923810	802.485714	9.192517	2.912925
1.0				
std	9.135373	403.509100	8.106864	1.024165
0.0				
min	18.000000	102.000000	1.000000	1.000000
1.0				
25%	30.000000	465.000000	2.000000	2.000000
1.0				
50%	36.000000	802.000000	7.000000	3.000000
1.0				
75%	43.000000	1157.000000	14.000000	4.000000
1.0				
max	60.000000	1499.000000	29.000000	5.000000
1.0				

EmployeeNumber	EnvironmentSatisfaction	HourlyRate
JobInvolvement \		
count	1470.000000	1470.000000
1470.000000		
mean	1024.865306	2.721769
2.729932		
std	602.024335	1.093082
0.711561		
min	1.000000	1.000000
1.000000		
25%	491.250000	2.000000
2.000000		
50%	1020.500000	3.000000
3.000000		
75%	1555.750000	4.000000
3.000000		
max	2068.000000	4.000000
4.000000		

JobLevel	...	RelationshipSatisfaction	StandardHours \
----------	-----	--------------------------	-----------------

count	1470.000000	...	1470.000000	1470.0
mean	2.063946	...	2.712245	80.0
std	1.106940	...	1.081209	0.0
min	1.000000	...	1.000000	80.0
25%	1.000000	...	2.000000	80.0
50%	2.000000	...	3.000000	80.0
75%	3.000000	...	4.000000	80.0
max	5.000000	...	4.000000	80.0

	StockOptionLevel	TotalWorkingYears	TrainingTimesLastYear	\
count	1470.000000	1470.000000	1470.000000	
mean	0.793878	11.279592	2.799320	
std	0.852077	7.780782	1.289271	
min	0.000000	0.000000	0.000000	
25%	0.000000	6.000000	2.000000	
50%	1.000000	10.000000	3.000000	
75%	1.000000	15.000000	3.000000	
max	3.000000	40.000000	6.000000	

	WorkLifeBalance	YearsAtCompany	YearsInCurrentRole	\
count	1470.000000	1470.000000	1470.000000	
mean	2.761224	7.008163	4.229252	
std	0.706476	6.126525	3.623137	
min	1.000000	0.000000	0.000000	
25%	2.000000	3.000000	2.000000	
50%	3.000000	5.000000	3.000000	
75%	3.000000	9.000000	7.000000	
max	4.000000	40.000000	18.000000	

	YearsSinceLastPromotion	YearsWithCurrManager
count	1470.000000	1470.000000
mean	2.187755	4.123129
std	3.222430	3.568136
min	0.000000	0.000000
25%	0.000000	2.000000
50%	1.000000	3.000000
75%	3.000000	7.000000
max	15.000000	17.000000

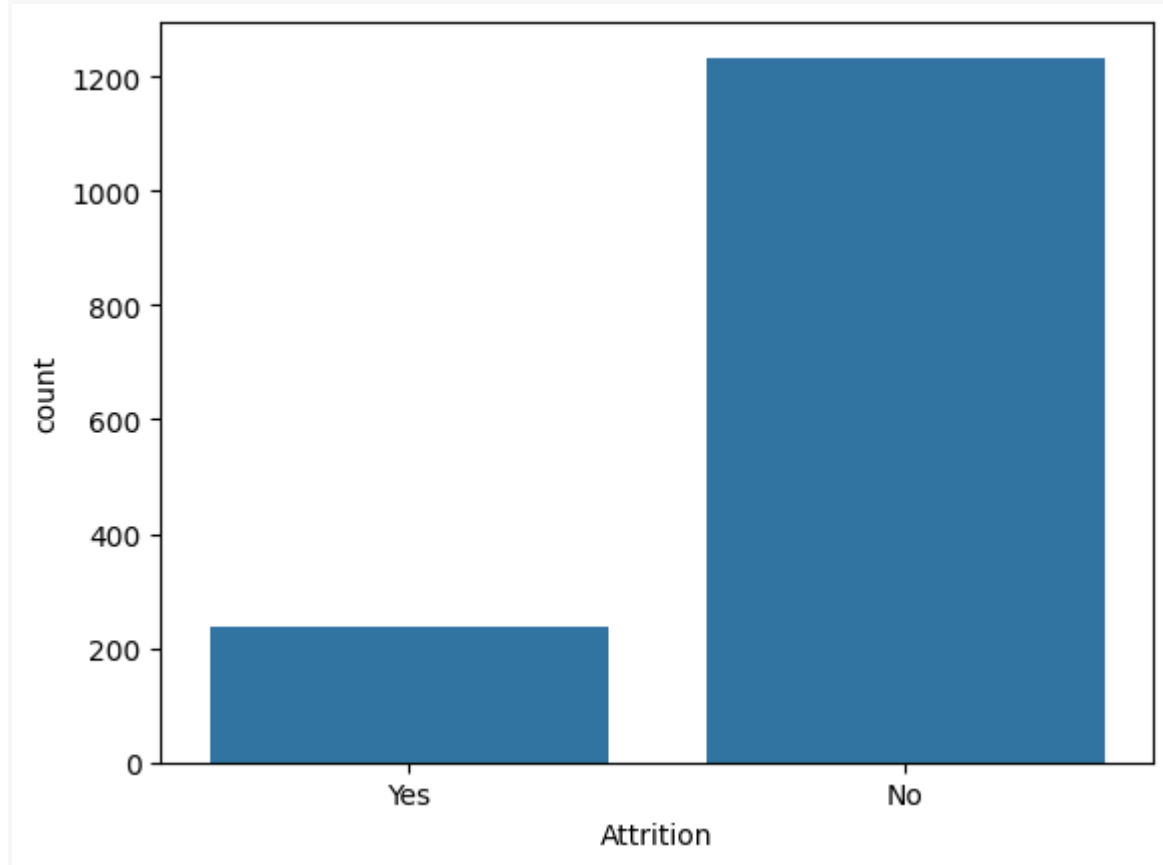
[8 rows x 26 columns]

4) Data Visualization

1. Co"→t plot roí tkc taígct :aíiable 'Attritio→t':

```
sns.countplot(x='Attrition', data=df)
```

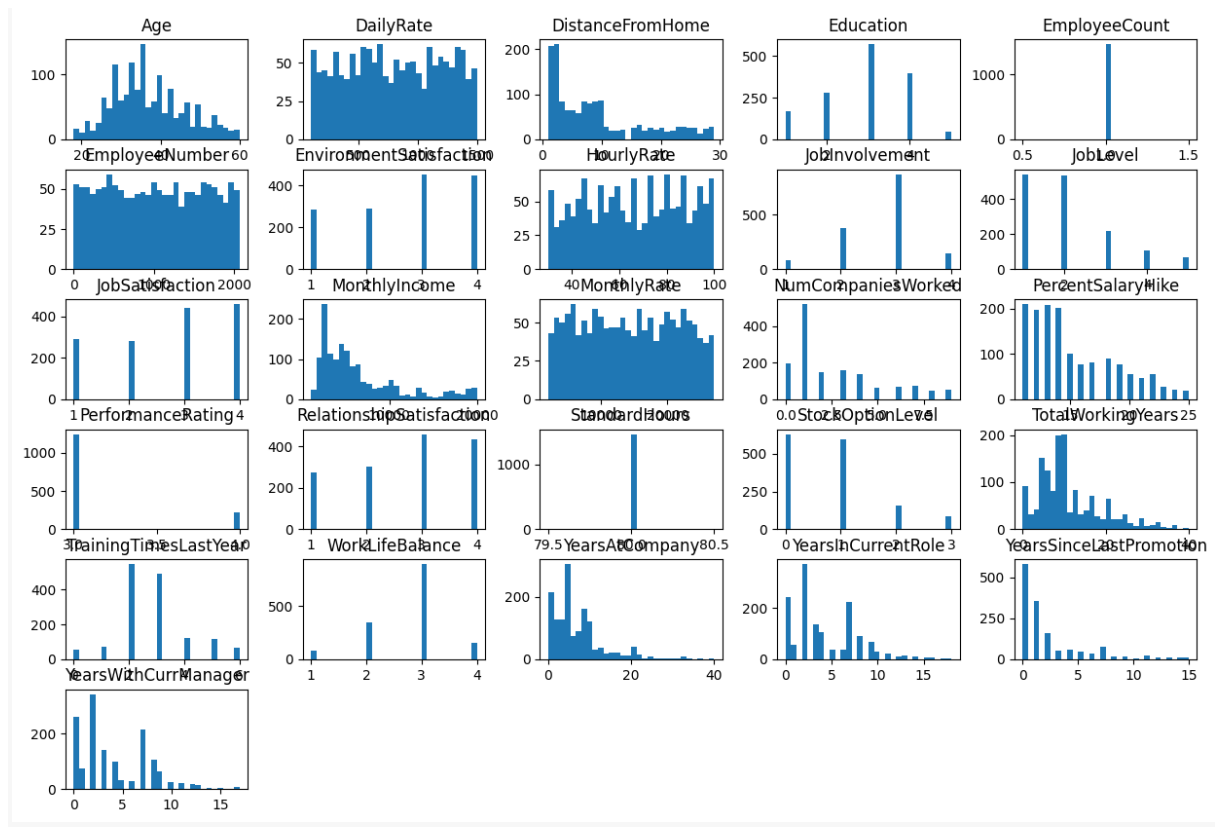
```
plt.show()
```



The count plot illustrates the distribution of employee attrition within the dataset. It is evident that the number of current employees ('No' attrition) significantly surpasses the number of employees who have left the company ('Yes' attrition). Such a distribution suggests that the dataset is imbalanced with respect to the target variable, which is an important characteristic to consider when developing predictive models, as it may influence model performance and will likely require specific techniques to handle the imbalance during the modeling phase.

2. Visualizaco da distribuio ou "anlise de catgorias"

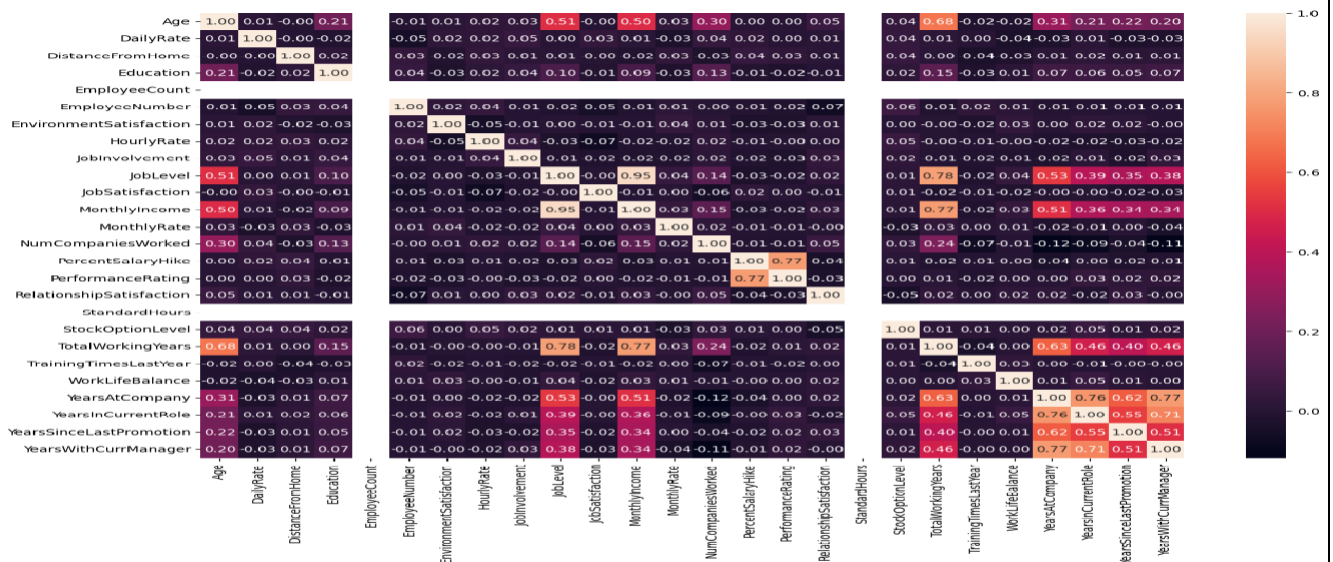
```
df.hist(bins=30, figsize=(15, 10), grid=False)  
plt.show()
```



The histogram plots reveal the underlying distributions of the numeric features in the dataset. Features related to employee satisfaction and ratings tend to show a multi-modal distribution, likely reflecting the discrete nature of survey responses. Salary-related features, such as `MonthlyIncome` and `PercentSalaryHike`, along with tenure-related features, like `TotalWorkingYears` and `YearsAtCompany`, exhibit right-skewed distributions. This skewness indicates that a smaller proportion of the workforce has very high salaries or has been with the company for an extended period. The uniform distributions of `DailyRate` and `HourlyRate` suggest variability in compensation that does not concentrate around a particular figure. Additionally, some features like `EmployeeCount` may not vary across the workforce, indicating they might not provide discriminative power in predictive modeling of employee attrition.

«Coíclatio» matíx kcatmap


```
plt.figure(figsize=(15, 10))
sns.heatmap(df.corr(), annot=True, fmt=".2f")
plt.show()
```



The correlation matrix heatmap provides a visual representation of the relationship strength between numerical features. The intensity of the colors corresponds to the magnitude of the correlation coefficient, where warmer colors denote higher positive correlations and cooler colors indicate negative correlations. Diagonal elements are maximally correlated as they represent the correlation of each variable with itself. Notable correlations such as those between TotalWorkingYears and JobLevel suggest a relationship where employees with more years of experience are in higher job positions. Such insights can be valuable for predictive modeling and hypothesis generation. It's also important to consider these correlations for feature selection to mitigate potential multicollinearity issues in machine learning models.

«.ExploíatolíQ Kata A→íalQsis


```
plt.figure(figsize=(6, 4))
sns.countplot(data=df, x='Gender', hue='Attrition')
plt.title('Gender vs Attrition')
plt.show()
```



The visualization provides a comparative view of the age distribution among current and former employees. The stacked histogram, complemented by the KDE curves, suggests a higher attrition rate among younger employees, while older employees show a higher retention rate. This pattern could indicate that age is an influential factor in employee turnover, with potential implications for HR policies and retention strategies. The density estimation curves help to highlight the differences in age distributions beyond the specific bin choices of the histogram, providing a smoother overview of the underlying age-related trends in attrition.

Gender vs Attrition


```
plt.figure(figsize=(6, 4))
sns.countplot(data=df, x='Gender', hue='Attrition')
plt.title('Gender vs Attrition')
plt.show()
```



The visualization provides a comparative view of the age distribution among current and former employees. The stacked histogram, complemented by the KDE curves, suggests a higher attrition rate among younger employees, while older employees show a higher retention rate. This pattern could indicate that age is an influential factor in employee turnover, with potential implications for HR policies and retention strategies. The density estimation curves help to highlight the differences in age distributions beyond the specific bin choices of the histogram, providing a smoother overview of the underlying age-related trends in attrition.

Department vs Attrition.

```
plt.figure(figsize=(8, 5))
sns.countplot(data=df, x='Department', hue='Attrition')
plt.title('Department vs Attrition')
plt.xticks(rotation=45)
plt.show()
```





The Department vs Attrition bar chart elucidates the distribution of employee attrition across different departments within the organization. The Research & Development department, being the largest, exhibits the highest counts of both retained and departed employees. Sales, while smaller in size, also shows a considerable attrition count. Human Resources has the lowest overall count, consistent with its smaller department size. When examining the proportion of attrition, it is crucial to consider the relative department sizes as well as the absolute counts to gain a true understanding of attrition patterns within each department.

Job vs Attrition


```
plt.figure(figsize=(12, 7))
sns.countplot(data=df, x='JobRole', hue='Attrition')
plt.title('Job Role vs Attrition')
plt.xticks(rotation=90)
plt.show()
```

The Job Role vs Attrition bar chart delineates the attrition rate across various job roles within the company. While roles like Sales Executive and Research Scientist have a higher absolute number of employees staying and leaving, the Laboratory Technician role stands out with a relatively high attrition rate compared to its size. In contrast, leadership roles such as Managers and Research Directors exhibit lower attrition rates, which aligns with expectations that higher job roles tend to have greater job stability. The data suggests that attrition is not uniformly distributed across job roles, and certain positions may require more attention to understand and address the underlying causes of turnover.



WorkLifeBalance vs Attrition

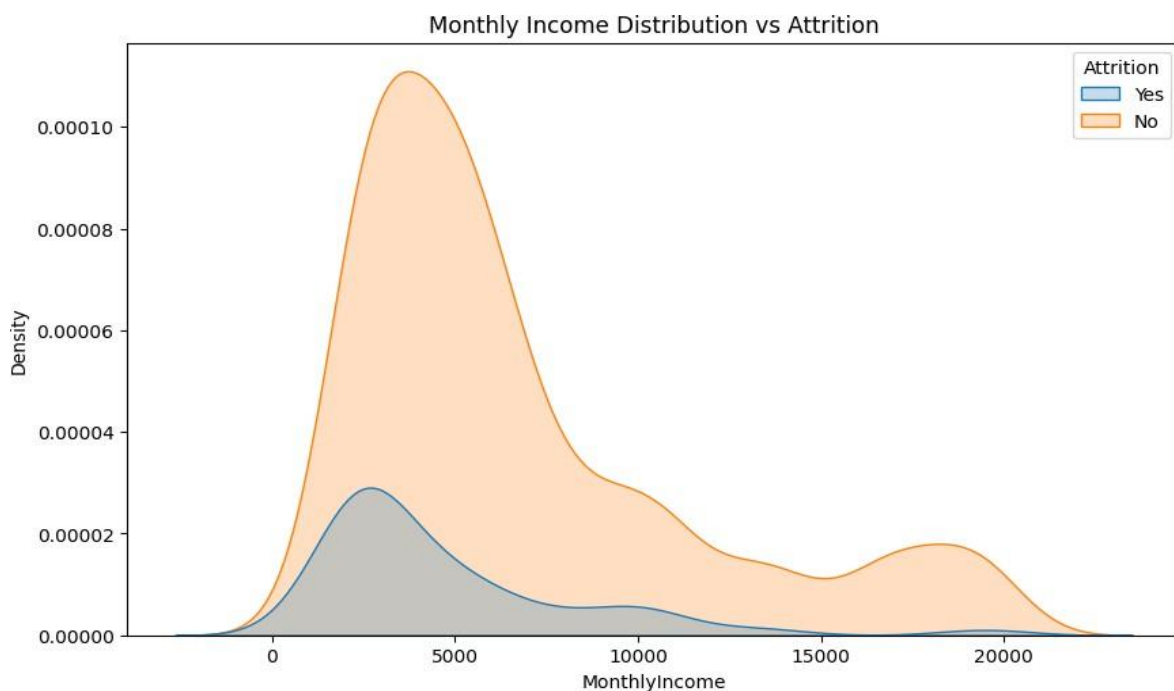
```
plt.figure(figsize=(8, 5))
sns.countplot(data=df, x='WorkLifeBalance', hue='Attrition')
plt.title('WorkLifeBalance vs Attrition')
plt.show()
```



The WorkLifeBalance vs Attrition chart depicts the correlation between employees' contentment with work-life balance and their decision to stay with or leave the company. A discernible trend shows that employees who report lower work-life balance ratings tend to leave the company at higher rates. Conversely, the highest work-life balance rating (4) corresponds with the lowest attrition, suggesting that employees who are most satisfied with their work-life integration are more inclined to remain at the company. This pattern underscores the importance of work-life balance as a factor in employee retention strategies.

Monthly Income Distribution vs Attrition.


```
plt.figure(figsize=(10, 6))
sns.kdeplot(data=df, x='MonthlyIncome', hue='Attrition', fill=True)
plt.title('Monthly Income Distribution vs Attrition')
plt.show()
```



The Monthly Income Distribution vs Attrition KDE plot provides insight into the role of compensation in employee turnover. The plot reveals that employees with lower monthly incomes are more densely represented among those who have left the company, suggesting a trend where lower income may contribute to higher turnover rates. In contrast, the distribution for employees who remain with the company extends towards higher income levels, which could indicate that competitive compensation is effective for employee retention. This pattern highlights the importance of considering compensation strategies when addressing workforce stability concerns.

Years at Company vs Attrition.

```
plt.figure(figsize=(10, 6))
sns.histplot(data=df, x='YearsAtCompany', hue='Attrition',
multiple='stack', kde=True)
plt.title('Years at Company vs Attrition')
plt.show()
```



The Years at Company vs Attrition chart provides a clear depiction of tenure's relation to employee attrition. Notably, there is a higher frequency of attrition among employees with shorter tenures, as demonstrated by the early peak and quick tapering off of the 'Yes' distribution. In contrast, the 'No' distribution is more spread out, reflecting that employees who stay with the company are likely to do so over many years, with a substantial number having long-term tenures. This pattern highlights the potential of tenure as a predictor of employee retention and may point to the benefits of developing strategies that encourage long-term commitment from the workforce.

Preparing a test set for the model

```
# Assuming 'Attrition' is the target variable and it's binary (Yes/No)
df['Attrition'] = df['Attrition'].apply(lambda x: 1 if x == 'Yes' else
0)

# Splitting dataset into features (X) and target (y)
X = df.drop('Attrition', axis=1)
y = df['Attrition']
```

```

# Encoding categorical variables and scaling numerical variables
categorical_features = [col for col in X.columns if X[col].dtype ==
'object']
numerical_features = [col for col in X.columns if col not in
categorical_features]

preprocessor = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), numerical_features),
        ('cat', OneHotEncoder(), categorical_features)
    ])

# Splitting the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
models = {
    "Logistic Regression": LogisticRegression(max_iter=1000),
    "Decision Tree": DecisionTreeClassifier(),
    "Random Forest": RandomForestClassifier(),
    "Gradient Boosting": GradientBoostingClassifier(),
    "SVM": SVC(probability=True),
    "KNN": KNeighborsClassifier(),
    "Neural Network": MLPClassifier(max_iter=1000),
    "XGBoost": XGBClassifier(use_label_encoder=False,
eval_metric='logloss')
}

results = {}

for name, model in models.items():
    # Pipeline for preprocessing and model training
    pipeline = Pipeline(steps=[('preprocessor', preprocessor),
                              ('model', model)])

    pipeline.fit(X_train, y_train)
    predictions = pipeline.predict(X_test)
    accuracy = accuracy_score(y_test, predictions)

    results[name] = {
        'Accuracy': accuracy,
        'Precision': precision_score(y_test, predictions),
        'Recall': recall_score(y_test, predictions),

```

```

        'F1 Score': f1_score(y_test, predictions)
    }

# Displaying performance metrics for each model
for result in results:
    print(f"{result}: {results[result]}")

```

Result:

```

Logistic Regression: {'Accuracy': 0.8945578231292517, 'Precision':
0.6428571428571429, 'Recall': 0.46153846153846156, 'F1 Score':
0.537313432835821}
Decision Tree: {'Accuracy': 0.7721088435374149, 'Precision':
0.18181818181818182, 'Recall': 0.20512820512820512, 'F1 Score':
0.1927710843373494}
Random Forest: {'Accuracy': 0.8741496598639455, 'Precision':
0.6666666666666666, 'Recall': 0.10256410256410256, 'F1 Score':
0.17777777777777778}
Gradient Boosting: {'Accuracy': 0.8809523809523809, 'Precision':
0.6666666666666666, 'Recall': 0.20512820512820512, 'F1 Score':
0.31372549019607837}
SVM: {'Accuracy': 0.891156462585034, 'Precision': 1.0, 'Recall':
0.1794871794871795, 'F1 Score': 0.30434782608695654}
KNN: {'Accuracy': 0.8639455782312925, 'Precision': 0.4444444444444444,
'Recall': 0.10256410256410256, 'F1 Score': 0.16666666666666666}
Neural Network: {'Accuracy': 0.8639455782312925, 'Precision':
0.4838709677419355, 'Recall': 0.38461538461538464, 'F1 Score':
0.4285714285714286}
XGBoost: {'Accuracy': 0.8775510204081632, 'Precision':
0.5882352941176471, 'Recall': 0.2564102564102564, 'F1 Score':
0.35714285714285715}

```

Example for tuning a Random Forest Classifier

```

param_grid = {
    'model__n_estimators': [100, 200],
    'model__max_depth': [10, 20, None],
    'model__min_samples_split': [2, 5],
    'model__min_samples_leaf': [1, 2]
}

pipeline = Pipeline(steps=[('preprocessor', preprocessor),
                           ('model', RandomForestClassifier())])

grid_search = GridSearchCV(pipeline, param_grid, cv=5,
scoring='accuracy')
grid_search.fit(X_train, y_train)

print("Best parameters:", grid_search.best_params_)
print("Best accuracy:", grid_search.best_score_)

```

Best parameters: {'model__max_depth': 10, 'model__min_samples_leaf': 1, 'model__min_samples_split': 2, 'model__n_estimators': 100}

Best accuracy: 0.8613883880274071

Train the final model (example with Random Forest)

```
final_model = grid_search.best_estimator_  
final_predictions = final_model.predict(X_test)
```

Final evaluation

```
print(f"Final Model Accuracy: {accuracy_score(y_test,  
final_predictions)}")  
print(classification_report(y_test, final_predictions))
```

Final Model Accuracy: 0.8707482993197279

	precision	recall	f1-score	support
0	0.88	0.99	0.93	255
1	0.57	0.10	0.17	39
accuracy			0.87	294
macro avg	0.72	0.55	0.55	294
weighted avg	0.84	0.87	0.83	294

Backend Flask Code :

```
import numpy as np  
import scipy as sp  
import pandas as pd  
from flask import Flask,request,jsonify,render_template  
import pickle  
  
app = Flask (__name__)  
model = pickle.load (open ('C:/Users/vaish/OneDrive/Documents/internship/model.pkl','rb'))  
  
@app.route('/')  
def home():  
    return render_template ('index.html')  
  
@app.route ('/predict',methods=['POST','GET'])  
def predict():  
    """  
    For rendering results on HTML GUI
```

```
"""
```

```
Age = request.form.get ("Age")
BusinessTravel = request.form['BusinessTravel']
DailyRate = request.form.get ('DailyRate')
Department = request.form['Department']
DistanceFromHome = request.form.get ("DistanceFromHome")
Education = request.form.get ("Education")
EducationField = request.form['EducationField']
EnvironmentSatisfaction = request.form.get ("EnvironmentSatisfaction")
Gender = request.form['Gender']
HourlyRate = request.form.get ("HourlyRate")
JobInvolvement = request.form.get ("EnvironmentSatisfaction")
JobLevel = request.form.get ("JobLevel")
JobRole = request.form['JobRole']
JobSatisfaction = request.form.get ("JobSatisfaction")
MaritalStatus = request.form['MaritalStatus']
MonthlyIncome = request.form.get ("MonthlyIncome")
NumCompaniesWorked = request.form.get ("NumCompaniesWorked")
OverTime = request.form['OverTime']
PerformanceRating = request.form.get ("PerformanceRating")
RelationshipSatisfaction = request.form.get ("RelationshipSatisfaction")
StockOptionLevel = request.form.get ("StockOptionLevel")
TotalWorkingYears = request.form.get ("TotalWorkingYears")
TrainingTimesLastYear = request.form.get ("TrainingTimesLastYear")
WorkLifeBalance = request.form.get ("WorkLifeBalance")
YearsAtCompany = request.form.get ("YearsAtCompany")
YearsInCurrentRole = request.form.get ("YearsInCurrentRole")
YearsSinceLastPromotion = request.form.get ("YearsSinceLastPromotion")
YearsWithCurrManager = request.form.get ("YearsWithCurrManager")
```

```
dict = {
    'Age': int (Age),
    'BusinessTravel': str (BusinessTravel),
    'DailyRate': int (DailyRate),
    'Department': Department,
    'DistanceFromHome': int (DistanceFromHome),
    'Education': Education,
    'EducationField': str (EducationField),
    'EnvironmentSatisfaction': int (EnvironmentSatisfaction),
    'Gender': str (Gender),
    'HourlyRate': int (HourlyRate),
    'JobInvolvement': int (JobInvolvement),
    'JobLevel': int (JobLevel),
    'JobRole': JobRole,
    'JobSatisfaction': int (JobSatisfaction),
    'MaritalStatus': str (MaritalStatus),
    'MonthlyIncome': int (MonthlyIncome),
    'NumCompaniesWorked': int (NumCompaniesWorked),
    'OverTime': str (OverTime),
    'PerformanceRating': int (PerformanceRating),
    'RelationshipSatisfaction': int (RelationshipSatisfaction),
```



```

        'StockOptionLevel': StockOptionLevel,
        'TotalWorkingYears': int (TotalWorkingYears),
        'TrainingTimesLastYear': TrainingTimesLastYear,
        'WorkLifeBalance': int (WorkLifeBalance),
        'YearsAtCompany': int (YearsAtCompany),
        'YearsInCurrentRole': int (YearsInCurrentRole),
        'YearsSinceLastPromotion': int (YearsSinceLastPromotion),
        'YearsWithCurrManager': int (YearsWithCurrManager)
    }

df = pd.DataFrame ([dict])

df['Total_Satisfaction'] = (df['EnvironmentSatisfaction'] +
                           df['JobInvolvement'] +
                           df['JobSatisfaction'] +
                           df['RelationshipSatisfaction'] +
                           df['WorkLifeBalance']) / 5

# Drop Columns
df.drop (
    ['EnvironmentSatisfaction', 'JobInvolvement', 'JobSatisfaction', 'RelationshipSatisfaction', 'WorkLifeBalance'],
    axis=1, inplace=True)

# Convert Total satisfaction into boolean
df['Total_Satisfaction_bool'] = df['Total_Satisfaction'].apply (lambda x: 1 if x >=
2.8 else 0)
df.drop ('Total_Satisfaction', axis=1, inplace=True)

# It can be observed that the rate of attrition of employees below age of 35 is high
df['Age_bool'] = df['Age'].apply (lambda x: 1 if x < 35 else 0)
df.drop ('Age', axis=1, inplace=True)

# It can be observed that the employees are more likely to drop the job if dailyRate
less than 800
df['DailyRate_bool'] = df['DailyRate'].apply (lambda x: 1 if x < 800 else 0)
df.drop ('DailyRate', axis=1, inplace=True)

# Employees working at R&D Department have higher attrition rate
df['Department_bool'] = df['Department'].apply (lambda x: 1 if x == 'Research &
Development' else 0)
df.drop ('Department', axis=1, inplace=True)

# Rate of attrition of employees is high if DistanceFromHome > 10
df['DistanceFromHome_bool'] = df['DistanceFromHome'].apply (lambda x: 1 if x > 10 else
0)
df.drop ('DistanceFromHome', axis=1, inplace=True)

# Employees are more likely to drop the job if the employee is working as Laboratory
Technician

```

```

df['JobRole_bool'] = df['JobRole'].apply (lambda x: 1 if x == 'Laboratory Technician'
else 0)
df.drop ('JobRole',axis=1,inplace=True)

# Employees are more likely to drop the job if the employee's hourly rate < 65
df['HourlyRate_bool'] = df['HourlyRate'].apply (lambda x: 1 if x < 65 else 0)
df.drop ('HourlyRate',axis=1,inplace=True)

# Employees are more likely to drop the job if the employee's MonthlyIncome < 4000
df['MonthlyIncome_bool'] = df['MonthlyIncome'].apply (lambda x: 1 if x < 4000 else 0)
df.drop ('MonthlyIncome',axis=1,inplace=True)

# Rate of attrition of employees is high if NumCompaniesWorked < 3
df['NumCompaniesWorked_bool'] = df['NumCompaniesWorked'].apply (lambda x: 1 if x > 3
else 0)
df.drop ('NumCompaniesWorked',axis=1,inplace=True)

# Employees are more likely to drop the job if the employee's TotalWorkingYears < 8
df['TotalWorkingYears_bool'] = df['TotalWorkingYears'].apply (lambda x: 1 if x < 8
else 0)
df.drop ('TotalWorkingYears',axis=1,inplace=True)

# Employees are more likely to drop the job if the employee's YearsAtCompany < 3
df['YearsAtCompany_bool'] = df['YearsAtCompany'].apply (lambda x: 1 if x < 3 else 0)
df.drop ('YearsAtCompany',axis=1,inplace=True)

# Employees are more likely to drop the job if the employee's YearsInCurrentRole <
3
df['YearsInCurrentRole_bool'] = df['YearsInCurrentRole'].apply (lambda x: 1 if x < 3
else 0)
df.drop ('YearsInCurrentRole',axis=1,inplace=True)

# Employees are more likely to drop the job if the employee's
YearsSinceLastPromotion < 1
df['YearsSinceLastPromotion_bool'] = df['YearsSinceLastPromotion'].apply (lambda x: 1
if x < 1 else 0)
df.drop ('YearsSinceLastPromotion',axis=1,inplace=True)

# Employees are more likely to drop the job if the employee's YearsWithCurrManager
< 1
df['YearsWithCurrManager_bool'] = df['YearsWithCurrManager'].apply (lambda x: 1 if x <
1 else 0)
df.drop ('YearsWithCurrManager',axis=1,inplace=True)

# Convert Categorical to Numerical
# Business Travel
if BusinessTravel == 'Rarely':
    df['BusinessTravel_Rarely'] = 1
    df['BusinessTravel_Frequently'] = 0
    df['BusinessTravel_No_Travel'] = 0
elif BusinessTravel == 'Frequently':

```

```

df['BusinessTravel_Rarely'] = 0
df['BusinessTravel_Frequently'] = 1
df['BusinessTravel_No_Travel'] = 0
else:
    df['BusinessTravel_Rarely'] = 0
    df['BusinessTravel_Frequently'] = 0
    df['BusinessTravel_No_Travel'] = 1
df.drop ('BusinessTravel',axis=1,inplace=True)

# Education
if Education == 1:
    df['Education_1'] = 1
    df['Education_2'] = 0
    df['Education_3'] = 0
    df['Education_4'] = 0
    df['Education_5'] = 0
elif Education == 2:
    df['Education_1'] = 0
    df['Education_2'] = 1
    df['Education_3'] = 0
    df['Education_4'] = 0
    df['Education_5'] = 0
elif Education == 3:
    df['Education_1'] = 0
    df['Education_2'] = 0
    df['Education_3'] = 1
    df['Education_4'] = 0
    df['Education_5'] = 0
elif Education == 4:
    df['Education_1'] = 0
    df['Education_2'] = 0
    df['Education_3'] = 0
    df['Education_4'] = 1
    df['Education_5'] = 0
else:
    df['Education_1'] = 0
    df['Education_2'] = 0
    df['Education_3'] = 0
    df['Education_4'] = 0
    df['Education_5'] = 1
df.drop ('Education',axis=1,inplace=True)

# EducationField
if EducationField == 'Life Sciences':
    df['EducationField_Life_Sciences'] = 1
    df['EducationField_Medical'] = 0
    df['EducationField_Marketing'] = 0
    df['EducationField_Technical_Degree'] = 0
    df['Education_Human_Resources'] = 0
    df['Education_Other'] = 0
elif EducationField == 'Medical':

```

```

df['EducationField_Life_Sciences'] = 0
df['EducationField_Medical'] = 1
df['EducationField_Marketing'] = 0
df['EducationField_Technical_Degree'] = 0
df['Education_Human_Resources'] = 0
df['Education_Other'] = 0
elif EducationField == 'Marketing':
    df['EducationField_Life_Sciences'] = 0
    df['EducationField_Medical'] = 0
    df['EducationField_Marketing'] = 1
    df['EducationField_Technical_Degree'] = 0
    df['Education_Human_Resources'] = 0
    df['Education_Other'] = 0
elif EducationField == 'Technical Degree':
    df['EducationField_Life_Sciences'] = 0
    df['EducationField_Medical'] = 0
    df['EducationField_Marketing'] = 0
    df['EducationField_Technical_Degree'] = 1
    df['Education_Human_Resources'] = 0
    df['Education_Other'] = 0
elif EducationField == 'Human Resources':
    df['EducationField_Life_Sciences'] = 0
    df['EducationField_Medical'] = 0
    df['EducationField_Marketing'] = 0
    df['EducationField_Technical_Degree'] = 0
    df['Education_Human_Resources'] = 1
    df['Education_Other'] = 0
else:
    df['EducationField_Life_Sciences'] = 0
    df['EducationField_Medical'] = 0
    df['EducationField_Marketing'] = 0
    df['EducationField_Technical_Degree'] = 0
    df['Education_Human_Resources'] = 1
    df['Education_Other'] = 1
df.drop ('EducationField',axis=1,inplace=True)

# Gender
if Gender == 'Male':
    df['Gender_Male'] = 1
    df['Gender_Female'] = 0
else:
    df['Gender_Male'] = 0
    df['Gender_Female'] = 1
df.drop ('Gender',axis=1,inplace=True)

# Marital Status
if MaritalStatus == 'Married':
    df['MaritalStatus_Married'] = 1
    df['MaritalStatus_Single'] = 0
    df['MaritalStatus_Divorced'] = 0
elif MaritalStatus == 'Single':

```

```

df['MaritalStatus_Married'] = 0
df['MaritalStatus_Single'] = 1
df['MaritalStatus_Divorced'] = 0
else:
    df['MaritalStatus_Married'] = 0
    df['MaritalStatus_Single'] = 0
    df['MaritalStatus_Divorced'] = 1
df.drop ('MaritalStatus',axis=1,inplace=True)

# Overtime
if OverTime == 'Yes':
    df['OverTime_Yes'] = 1
    df['OverTime_No'] = 0
else:
    df['OverTime_Yes'] = 0
    df['OverTime_No'] = 1
df.drop ('OverTime',axis=1,inplace=True)

# Stock Option Level
if StockOptionLevel == 0:
    df['StockOptionLevel_0'] = 1
    df['StockOptionLevel_1'] = 0
    df['StockOptionLevel_2'] = 0
    df['StockOptionLevel_3'] = 0
elif StockOptionLevel == 1:
    df['StockOptionLevel_0'] = 0
    df['StockOptionLevel_1'] = 1
    df['StockOptionLevel_2'] = 0
    df['StockOptionLevel_3'] = 0
elif StockOptionLevel == 2:
    df['StockOptionLevel_0'] = 0
    df['StockOptionLevel_1'] = 0
    df['StockOptionLevel_2'] = 1
    df['StockOptionLevel_3'] = 0
else:
    df['StockOptionLevel_0'] = 0
    df['StockOptionLevel_1'] = 0
    df['StockOptionLevel_2'] = 0
    df['StockOptionLevel_3'] = 1
df.drop ('StockOptionLevel',axis=1,inplace=True)

# Training Time Last Year
if TrainingTimesLastYear == 0:
    df['TrainingTimesLastYear_0'] = 1
    df['TrainingTimesLastYear_1'] = 0
    df['TrainingTimesLastYear_2'] = 0
    df['TrainingTimesLastYear_3'] = 0
    df['TrainingTimesLastYear_4'] = 0
    df['TrainingTimesLastYear_5'] = 0
    df['TrainingTimesLastYear_6'] = 0
elif TrainingTimesLastYear == 1:

```

```

df['TrainingTimesLastYear_0'] = 0
df['TrainingTimesLastYear_1'] = 1
df['TrainingTimesLastYear_2'] = 0
df['TrainingTimesLastYear_3'] = 0
df['TrainingTimesLastYear_4'] = 0
df['TrainingTimesLastYear_5'] = 0
df['TrainingTimesLastYear_6'] = 0
elif TrainingTimesLastYear == 2:
    df['TrainingTimesLastYear_0'] = 0
    df['TrainingTimesLastYear_1'] = 0
    df['TrainingTimesLastYear_2'] = 1
    df['TrainingTimesLastYear_3'] = 0
    df['TrainingTimesLastYear_4'] = 0
    df['TrainingTimesLastYear_5'] = 0
    df['TrainingTimesLastYear_6'] = 0
elif TrainingTimesLastYear == 3:
    df['TrainingTimesLastYear_0'] = 0
    df['TrainingTimesLastYear_1'] = 0
    df['TrainingTimesLastYear_2'] = 0
    df['TrainingTimesLastYear_3'] = 1
    df['TrainingTimesLastYear_4'] = 0
    df['TrainingTimesLastYear_5'] = 0
    df['TrainingTimesLastYear_6'] = 0
elif TrainingTimesLastYear == 4:
    df['TrainingTimesLastYear_0'] = 0
    df['TrainingTimesLastYear_1'] = 0
    df['TrainingTimesLastYear_2'] = 0
    df['TrainingTimesLastYear_3'] = 0
    df['TrainingTimesLastYear_4'] = 1
    df['TrainingTimesLastYear_5'] = 0
    df['TrainingTimesLastYear_6'] = 0
elif TrainingTimesLastYear == 5:
    df['TrainingTimesLastYear_0'] = 0
    df['TrainingTimesLastYear_1'] = 0
    df['TrainingTimesLastYear_2'] = 0
    df['TrainingTimesLastYear_3'] = 0
    df['TrainingTimesLastYear_4'] = 0
    df['TrainingTimesLastYear_5'] = 1
    df['TrainingTimesLastYear_6'] = 0
else:
    df['TrainingTimesLastYear_0'] = 0
    df['TrainingTimesLastYear_1'] = 0
    df['TrainingTimesLastYear_2'] = 0
    df['TrainingTimesLastYear_3'] = 0
    df['TrainingTimesLastYear_4'] = 0
    df['TrainingTimesLastYear_5'] = 0
    df['TrainingTimesLastYear_6'] = 1
df.drop ('TrainingTimesLastYear',axis=1,inplace=True)

df.to_csv ('features.csv',index=False)

```

```
prediction = model.predict (df)

if prediction == 0:
    return render_template ('index.html',prediction_text='Employee Might Not Leave The
Job')

else:
    return render_template ('index.html',prediction_text='Employee Might Leave The
Job')

# Drop Columns

# Convert Total satisfaction into boolean

# Convert Categorical to Numerical
# Buisness Travel

# Education

# df.to_csv ('features.csv',index=False)
print(df)

if __name__ == "__main__":
    app.run (debug=True)
```

Basic Front-End UI :

Employee Attrition Prediction

Age:

Business Travel:

Daily Rate:

Department:

Distance From Home:

Education:

Education Field:

Environment Satisfaction:

Gender:

Hourly Rate:

Job Involvement:

Job Level:

Job Role:

Job Satisfaction:

Marital Status:

Marital Status:

Monthly Income:

Number of Companies Worked in:

Over Time:

Performance Rating:

Relationship Satisfaction:

Stock Option Level:

Total Working Years:

Training Times Last Year:

Work Life Balance:

Years At Company:

Years In Current Role:

Years Since Last Promotion:

Years with current manger:

Predict

Employee Might Leave The Job

CONCLUSION:

IN CONCLUSION, USING MACHINE LEARNING TO MANAGE EMPLOYEE ATTRITION IS A POWERFUL WAY FOR COMPANIES TO TACKLE TURNOVER ISSUES AND BUILD A STRONGER WORKFORCE. BY ANALYZING DATA AND PREDICTING WHICH EMPLOYEES MIGHT LEAVE, ORGANIZATIONS CAN TAKE PROACTIVE STEPS TO KEEP THEM ENGAGED AND SATISFIED. THIS INCLUDES THINGS LIKE PERSONALIZED RETENTION STRATEGIES, IMPROVING HOW THEY RECRUIT AND DEVELOP TALENT, AND CREATING A POSITIVE WORKPLACE CULTURE. NOT ONLY DOES THIS APPROACH HELP SAVE MONEY BY REDUCING THE COSTS OF TURNOVER, BUT IT ALSO FOSTERS A WORK ENVIRONMENT WHERE EMPLOYEES FEEL VALUED AND SUPPORTED. PLUS, BY MAKING SMARTER DECISIONS BASED ON DATA, COMPANIES CAN BETTER PLAN FOR THE FUTURE AND STAY COMPETITIVE IN THEIR INDUSTRY. OVERALL, EMBRACING MACHINE LEARNING IN EMPLOYEE ATTRITION MANAGEMENT LEADS TO HAPPIER EMPLOYEES, STRONGER TEAMS, AND GREATER SUCCESS FOR THE COMPANY.

